35_ml_101

March 23, 2022

1 Machine Learning 101

[ ]:
```
# importing libraries
import pandas as pd
import seaborn as sns
import matplotlib.pyplot as plt
from sklearn.model_selection import train_test_split
from sklearn.linear_model import LinearRegression
plt.style.use({'figure.facecolor':'white'})
```

[ ]:
```
data = pd.read_csv('../../data/salary_data.csv')
data.head()
```

[ ]:

	YearsExperience	Salary
0	1.1	39343
1	1.3	46205
2	1.5	37731
3	2.0	43525
4	2.2	39891

[ ]:
```
# X = data.iloc[:, :-1].values
# X
# type(X)
```

[ ]:
```
X = data.loc[:, ["YearsExperience"]].values
X
```

[ ]:
```
array([[ 1.1],
       [ 1.3],
       [ 1.5],
       [ 2. ],
       [ 2.2],
       [ 2.9],
       [ 3. ],
       [ 3.2],
       [ 3.2],
       [ 3.7],
       [ 3.9],
```

```
       [ 4. ],
       [ 4. ],
       [ 4.1],
       [ 4.5],
       [ 4.9],
       [ 5.1],
       [ 5.3],
       [ 5.9],
       [ 6. ],
       [ 6.8],
       [ 7.1],
       [ 7.9],
       [ 8.2],
       [ 8.7],
       [ 9. ],
       [ 9.5],
       [ 9.6],
       [10.3],
       [10.5]])
```

[ ]:
```
y = data.loc[:, ["Salary"]].values
y
```

[ ]:
```
array([[ 39343],
       [ 46205],
       [ 37731],
       [ 43525],
       [ 39891],
       [ 56642],
       [ 60150],
       [ 54445],
       [ 64445],
       [ 57189],
       [ 63218],
       [ 55794],
       [ 56957],
       [ 57081],
       [ 61111],
       [ 67938],
       [ 66029],
       [ 83088],
       [ 81363],
       [ 93940],
       [ 91738],
       [ 98273],
       [101302],
       [113812],
```

```
        [109431],
        [105582],
        [116969],
        [112635],
        [122391],
        [121872]], dtype=int64)
```

[ ]:
```
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.33,␣
↪random_state=1)
```

1.1 Exploring train and test data

[ ]:
```
X_train
```

[ ]:
```
array([[ 8.2],
       [ 2.2],
       [ 1.5],
       [ 9. ],
       [ 3. ],
       [ 5.9],
       [ 4.1],
       [ 3.2],
       [ 9.6],
       [ 1.3],
       [ 5.1],
       [ 1.1],
       [ 4.9],
       [10.5],
       [10.3],
       [ 3.7],
       [ 3.2],
       [ 4. ],
       [ 4. ],
       [ 2.9]])
```

[ ]:
```
X_test
```

[ ]:
```
array([[5.3],
       [7.1],
       [3.9],
       [6. ],
       [4.5],
       [6.8],
       [9.5],
       [2. ],
       [8.7],
       [7.9]])
```

[ ]:
```
y_train
```

[ ]:
```
array([[113812],
       [ 39891],
       [ 37731],
       [105582],
       [ 60150],
       [ 81363],
       [ 57081],
       [ 54445],
       [112635],
       [ 46205],
       [ 66029],
       [ 39343],
       [ 67938],
       [121872],
       [122391],
       [ 57189],
       [ 64445],
       [ 56957],
       [ 55794],
       [ 56642]], dtype=int64)
```

[ ]:
```
y_test
```

[ ]:
```
array([[ 83088],
       [ 98273],
       [ 63218],
       [ 93940],
       [ 61111],
       [ 91738],
       [116969],
       [ 43525],
       [109431],
       [101302]], dtype=int64)
```

[ ]:
```
regressor = LinearRegression()
regressor.fit(X=X_train, y=y_train)
```

[ ]:
```
LinearRegression()
```

[ ]:
```
help(regressor.fit)
```

[ ]:
```
# viz_train = plt
plt.scatter(X_train, y_train, color='red')
```

[ ]:
```
<matplotlib.collections.PathCollection at 0x2b1f54b9f60>
```

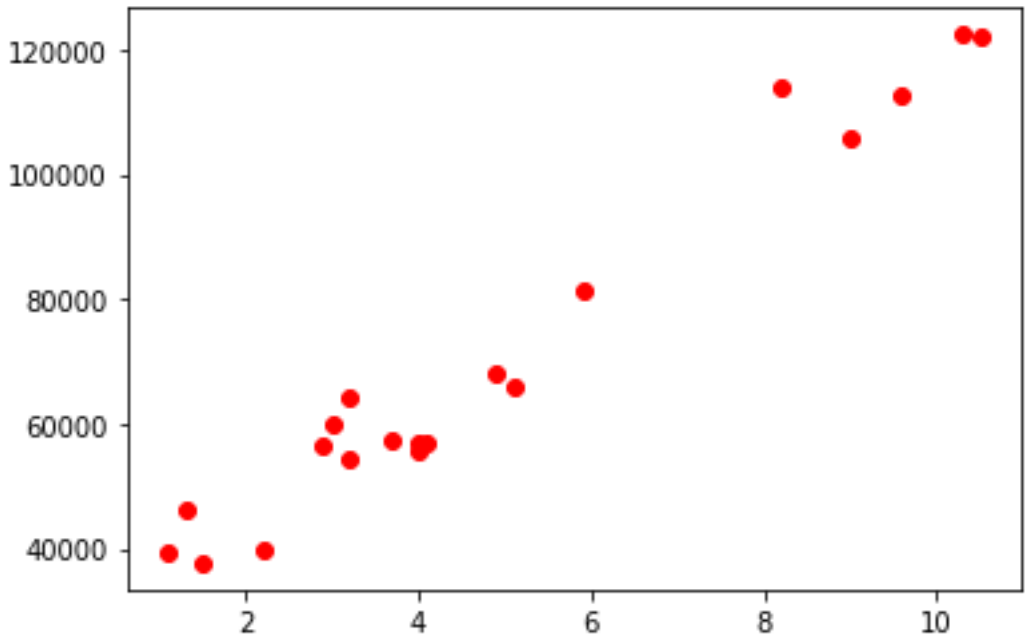

[ ]:
```
plt.plot(X_train, regressor.predict(X_train), color="blue")
```

[ ]:
```
[<matplotlib.lines.Line2D at 0x2b1f55d6230>]
```

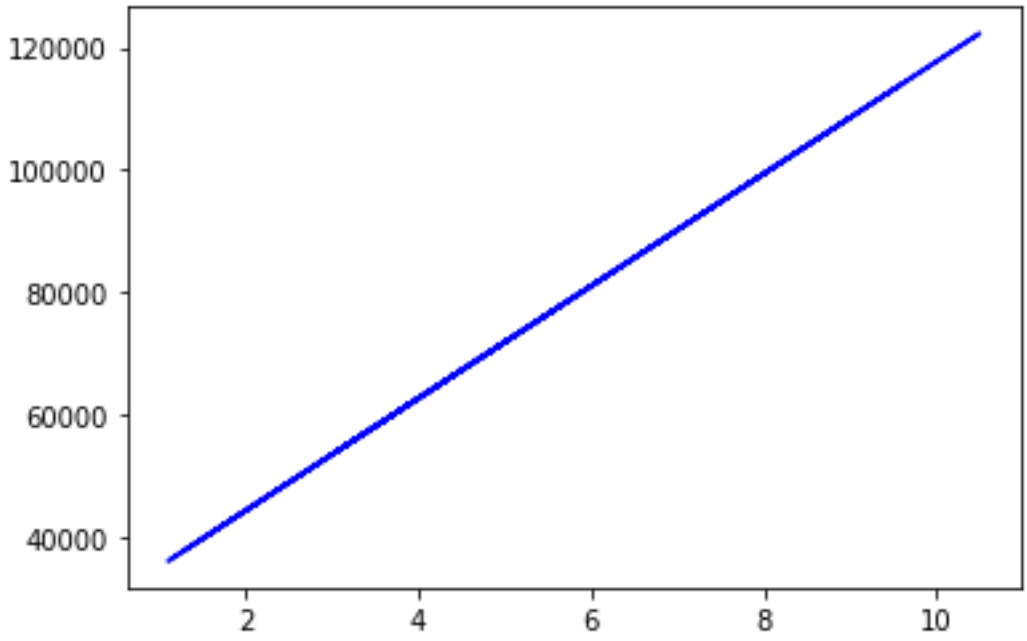

[ ]:
```
plt.scatter(X_train, y_train, color='red')
plt.plot(X_train, regressor.predict(X_train), color="blue")
```

[ ]:
```
[<matplotlib.lines.Line2D at 0x2b1f56303a0>]
```

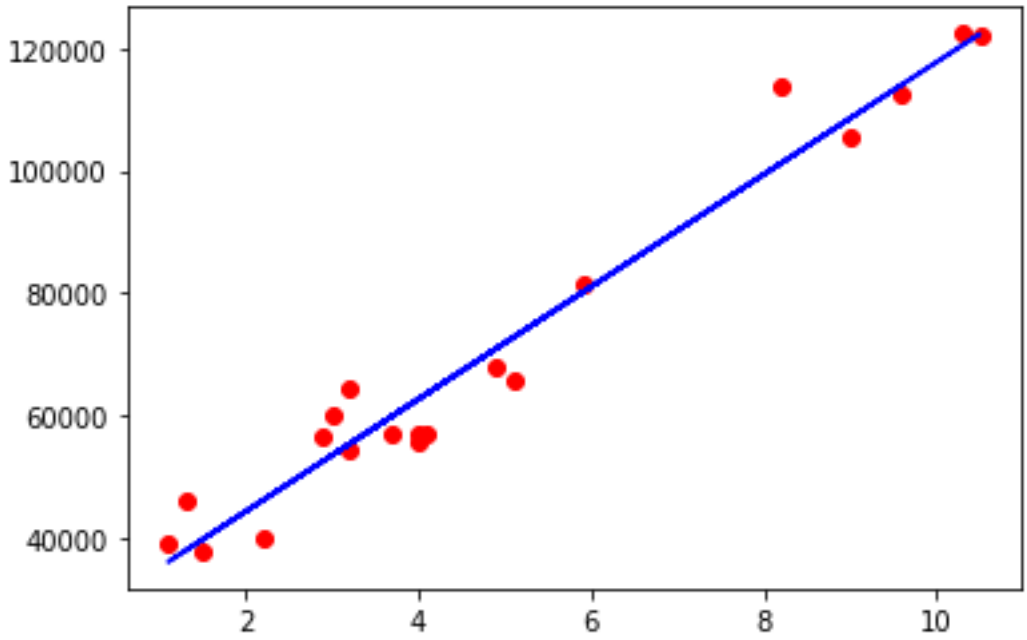

[ ]:
```
plt.scatter(X_train, y_train, color='red')
plt.plot(X_train, regressor.predict(X_train), color="blue")
plt.title('Salary VS Experience (Training Data Set)')
plt.xlabel('Year of Experience')
plt.ylabel('Salary')
plt.show()
```

[ ]:
```
plt.scatter(X_test, y_test, color='red')
plt.plot(X_test, regressor.predict(X_test), color="blue")
plt.title('Salary VS Experience (Test Data Set)')
plt.xlabel('Year of Experience')
plt.ylabel('Salary')
plt.show()
```

[ ]:
```
# predict 5 Years of experience's salary
y_pred_arr = regressor.predict(X_test)
y_pred_arr
```

[ ]:
```
array([[ 74675.37776747],
       [ 91160.02832519],
       [ 61853.98288925],
       [ 81086.07520659],
       [ 67348.86640849],
       [ 88412.58656557],
       [113139.56240215],
       [ 44453.51841166],
       [105813.05104316],
       [ 98486.53968418]])
```

[ ]:
```
y_pred = regressor.predict([[5]])
y_pred
```

[ ]:
```
array([[71927.93600785]])
```

[ ]:
```
data.corr()
```

[ ]:

	YearsExperience	Salary
YearsExperience	1.000000	0.978242
Salary	0.978242	1.000000

[ ]:
```
sns.heatmap(data.corr(), annot=True)
```

[ ]: <AxesSubplot:>